

# MENG 3065 - MODULE 6

## Artificial Intelligence: A Modern Approach Chapter 22 Reinforcement Learning

These slides has been extracted, modified and updated from original slides of :  
Artificial Intelligence: A Modern Approach, Fourth Edition.© 2022 Pearson Education, Inc.



**WE ARE**

**HUMBER**

# Outline

- An introduction to Reinforcement Learning
- Sequential Decision Problems
- Learning from Rewards
- Passive Reinforcement Learning
- Active Reinforcement Learning
- Examples of Reinforcement Learning
- Applications of Reinforcement Learning

# An introduction to Reinforcement Learning




<https://youtu.be/JgvyzlkxgF0?si=pa4KFoz7DaVOzmc3>

# Understanding Reinforcement Learning

- Distinct from Supervised Learning: No direct supervision with "correct" answers
- Scalar Reward Signal: Feedback provided through a numerical reward
- Temporal Considerations: Involves a notion of "time" in terms of steps or moves
- Delayed Feedback: Feedback is not instantaneous but occurs after actions are taken
- Dynamic Impact of Actions: Agent's actions influence subsequent data received

# Reinforcement Learning

- Environment
  - Provides the agent with the current state
  - Provides a reward at each time step
- Agent
  - Chooses an action at each time step, given the state
- Reward
  - A reward  $r$  is a scalar feedback signal
  - The goal of the agent is to select actions to maximize total future reward.



# Key Components of an RL Agent

- Policy:
  - Describes the agent's **behavioral strategy** or decision-making function
- Value Function:
  - Calculates the expected **future rewards** associated with a given state
- Model:
  - Represents the agent's internal **model** of the external **environment**

# Policy

- A policy defines the behavior of an agent by specifying the action it selects at a given state, often denoted by the symbol  $\pi$ .
- The policy function is a mapping that associates states with corresponding actions.
  - deterministic:  $a = \pi(s)$
  - stochastic:  $\pi(a | s) = P[A_t = a | S_t = s]$

# Value Function

- A value function anticipates future rewards under a specific policy.
- A tool to assess the goodness or badness of states.
- Maps the states to their expected discounted returns to help evaluate the quality of states.
- The Q-function extends this concept by mapping pairs (states, actions) to their corresponding expected discounted returns



# Model

- Predicts the next actions of the environment
  - predicts the next state given the current state and a specified action
  - predicts the next reward


# Sequential Decision Problems

- **Markov decision process (MDP):** a sequential decision problem for a fully observable, stochastic environment
- MDP consists of:
  - a set of states (with an initial state  $s_0$ );
  - a set  $\text{ACTIONS}(s)$  of actions in each state;
  - a transition model  $P(s' / s, a)$ ; and
  - a reward function  $R(s, a, s')$ .

# Sequential Decision Problems

- MDP solutions usually involve **dynamic programming** simplifying a problem by recursively breaking it into smaller pieces and remembering the optimal solutions to the pieces.
- A solution called **policy**.
  - specify what the agent should do for any state that the agent might reach
  - the quality of a policy is measured by the expected utility of possible environment histories generated
  - **optimal policy**: highest expected utility

# Sequential Decision Problems



a) A simple, stochastic 4x3 environment that presents the agent with a sequential decision problem.

(b) Illustration of the transition model of the environment: the “intended” outcome occurs with probability 0.8, but with probability 0.2 the agent moves at right angles to the intended direction. A collision with a wall results in no movement. Transitions into the two terminal states have reward +1 and -1, respectively, and all other transitions have a reward of -0.04.

# Sequential Decision Problems

- Utility of a state is the expected reward for the next transition plus the discounted utility of the next state, assuming that the agent chooses the optimal action

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')].$$

- This is called the **Bellman equation**, after Richard Bellman (1957).
- **Action-utility function, or Q-function:**  $Q(s, a)$ 
  - the expected utility of taking a given action in a given state.
  - related to utilities in the obvious way:

$$U(s) = \max_a Q(s, a).$$

- The optimal policy can be extracted from the Q-function

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

- The Q-function is in algorithms for solving MDPs

```
function Q-VALUE(mdp, s, a, U) returns a utility value
  return  $\sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U[s']]$ 
```

# Sequential Decision Problems

3	0.8516	0.9078	0.9578	+1
2	0.8016		0.7003	-1
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4

- The utilities of the states in the 4X3 world with  $\gamma = 1$  and  $r = -0.04$  for transitions to nonterminal states.
- In this example: states are defined as (col#, row#)
- The Bellman equation for the state (1, 1) is  $U(1,1) = \max \{r + \gamma * U(s')\}$

$$\max \{ \begin{array}{l} [0.8(-0.04 + \gamma U(1, 2)) + 0.1(-0.04 + \gamma U(2, 1)) + 0.1(-0.04 + \gamma U(1, 1))], \quad \leftarrow \text{up} \\ [0.9(-0.04 + \gamma U(1, 1)) + 0.1(-0.04 + \gamma U(1, 2))], \quad \leftarrow \text{left} \\ [0.9(-0.04 + \gamma U(1, 1)) + 0.1(-0.04 + \gamma U(2, 1))], \quad \leftarrow \text{down} \\ [0.8(-0.04 + \gamma U(2, 1)) + 0.1(-0.04 + \gamma U(1, 2)) + 0.1(-0.04 + \gamma U(1, 1))] \end{array} \} \quad \leftarrow \text{right}$$

# Learning from Rewards

- Agent interacts with the world and periodically receives rewards (reinforcements)
- Varieties of approaches:
- **Model-based reinforcement learning:** uses a transition model
  - Model may be initially unknown
  - Learns from observing effects of actions
  - Useful for state estimation
  - Learn a utility function  $U(s)$ ,
- **Model-free reinforcement learning:** neither knows nor learns transition model
  - **Action-utility learning:** most common form **Q-learning**, where the agent learns a **Q-function**, or quality-function,  $Q(s, a)$ , denoting the sum of rewards from state  $s$  if action  $a$  is taken.
- **Policy search:** learns a policy  $\pi(s)$  that maps directly from states to actions.

# Learning from Rewards

## Passive Reinforcement Learning

- **Passive learning agent:** agent that learns the utility function  $U^\pi(s)$ 
  - Expected total discounted reward if policy  $\pi$  is executed beginning in state  $s$
  - does not know the transition model  $P(s' | s, a)$ ,
  - executes a set of trials in the environment using its policy  $\pi$ . starts in state (1,1) and experiences a sequence of state transitions until it reaches one of the terminal states, (4,2) or (4,3).

(1,1)	$\xrightarrow[\text{Up}]{-.04}$	(1,2)	$\xrightarrow[\text{Up}]{-.04}$	(1,3)	$\xrightarrow[\text{Right}]{-.04}$	(1,2)	$\xrightarrow[\text{Up}]{-.04}$	(1,3)	$\xrightarrow[\text{Right}]{-.04}$	(2,3)	$\xrightarrow[\text{Right}]{-.04}$	(3,3)	$\xrightarrow[\text{Right}]{+1}$	(4,3)
(1,1)	$\xrightarrow[\text{Up}]{-.04}$	(1,2)	$\xrightarrow[\text{Up}]{-.04}$	(1,3)	$\xrightarrow[\text{Right}]{-.04}$	(2,3)	$\xrightarrow[\text{Right}]{-.04}$	(3,3)	$\xrightarrow[\text{Right}]{-.04}$	(3,2)	$\xrightarrow[\text{Up}]{-.04}$	(3,3)	$\xrightarrow[\text{Right}]{+1}$	(4,3)
(1,1)	$\xrightarrow[\text{Up}]{-.04}$	(1,2)	$\xrightarrow[\text{Up}]{-.04}$	(1,3)	$\xrightarrow[\text{Right}]{-.04}$	(2,3)	$\xrightarrow[\text{Right}]{-.04}$	(3,3)	$\xrightarrow[\text{Right}]{-.04}$	(3,2)	$\xrightarrow[\text{Up}]{-1}$	(4,2)		

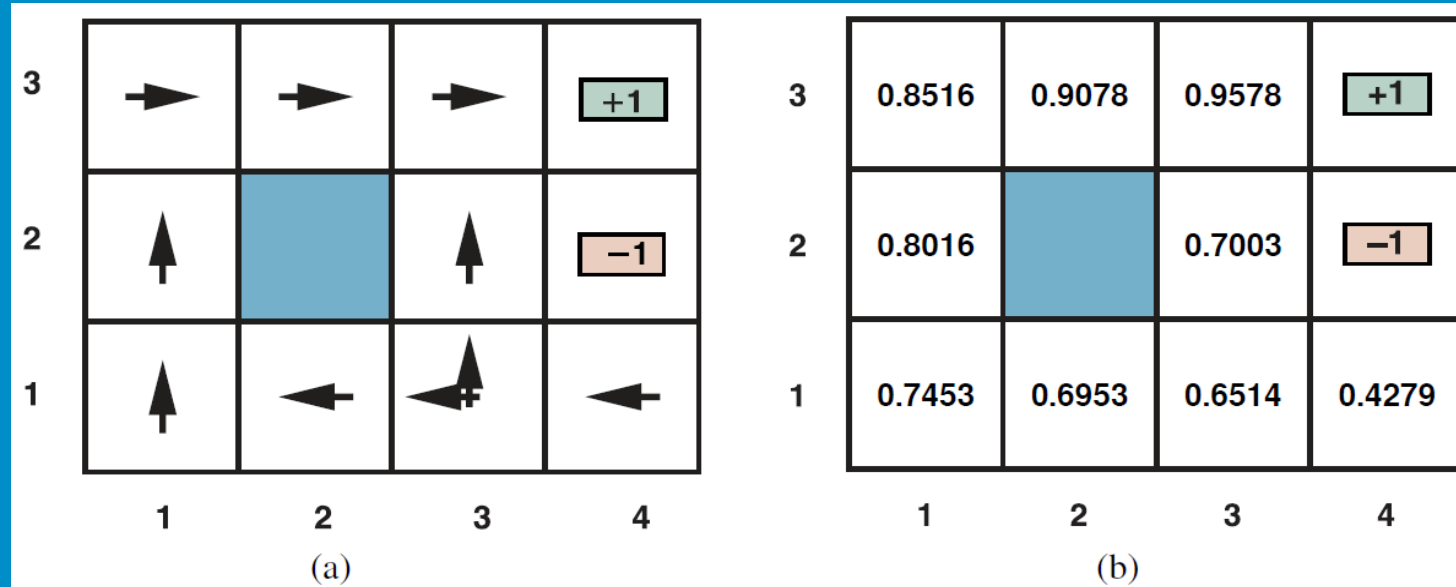
- Expected utility

$$U^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right],$$



# Learning from Rewards

## Passive Reinforcement Learning



- (a) The optimal policies for the stochastic environment with  $R(s, a, s^t) = 0.04$  for transitions between nonterminal states. There are two policies because in state (3,1) both *Left* and *Up* are optimal.
- (b) The utilities of the states in the  $4 \times 3$  world, given policy

# Learning from Rewards

## Direct utility estimation

- utility of a state is defined as the expected total reward from that state onward (reward-to-go)
- at the end of each sequence, the algorithm calculates the observed reward-to-go for each state and updates the estimated utility
- reduced reinforcement learning to a standard supervised learning problem in which each example is a *(state, reward-to-go)* pair.
- The utility of a state is determined by the reward and the expected utility of the successor states

$$U_i(s) = \sum_{s'} P(s' | s, \pi_i(s)) [R(s, \pi_i(s), s') + \gamma U_i(s')].$$

# Learning from Rewards

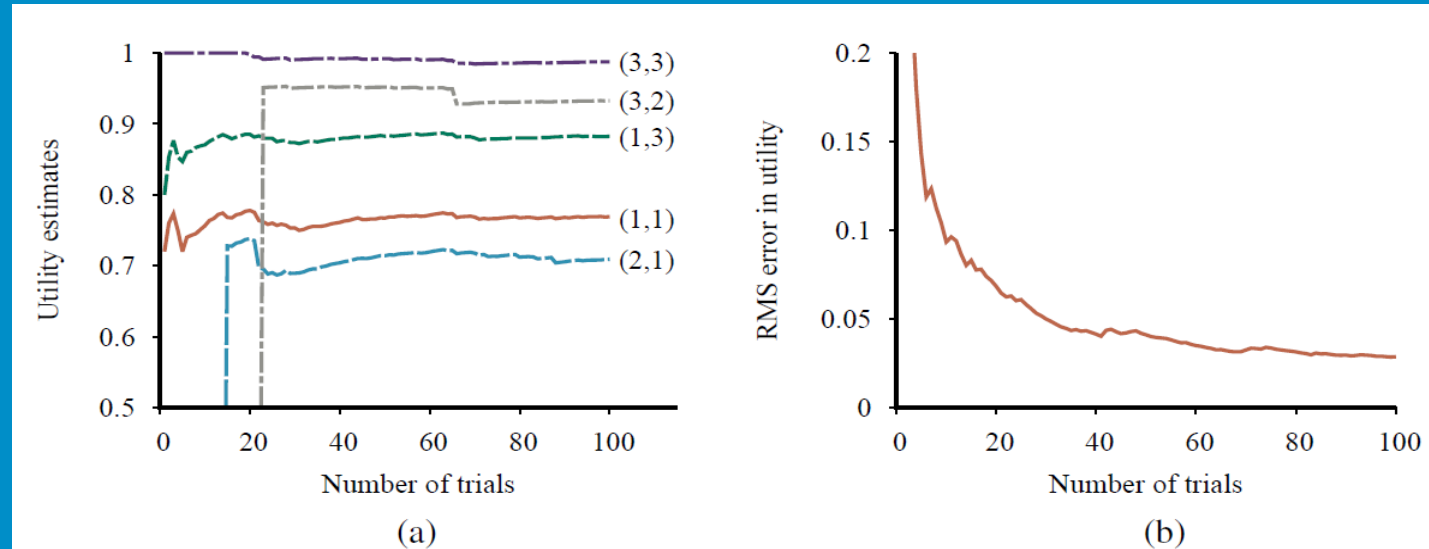
## Adaptive dynamic programming (ADP)

- Agent takes advantage of the constraints among the utilities of states by learning the transition model that connects them
- Solving the corresponding Markov decision process using dynamic programming

```
function PASSIVE-ADP-LEARNER(percept) returns an action
  inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r$ 
  persistent:  $\pi$ , a fixed policy
                $mdp$ , an MDP with model  $P$ , rewards  $R$ , actions  $A$ , discount  $\gamma$ 
                $U$ , a table of utilities for states, initially empty
                $N_{s'|s,a}$ , a table of outcome count vectors indexed by state and action, initially zero
                $s, a$ , the previous state and action, initially null

  if  $s'$  is new then  $U[s'] \leftarrow 0$ 
  if  $s$  is not null then
    increment  $N_{s'|s,a}[s, a][s']$ 
     $R[s, a, s'] \leftarrow r$ 
    add  $a$  to  $A[s]$ 
     $\mathbf{P}(\cdot \mid s, a) \leftarrow \text{NORMALIZE}(N_{s'|s,a}[s, a])$ 
     $U \leftarrow \text{POLICYEVALUATION}(\pi, U, mdp)$ 
     $s, a \leftarrow s', \pi[s']$ 
  return  $a$ 
```

# Learning from Rewards



The passive ADP learning curves for the 4 x 3 world.

- (a) The utility estimates for a selected subset of states, as a function of the number of trials. Notice that it takes 14 and 23 trials respectively before the rarely visited states (2,1) and (3,2) “discover” that they connect to the +1 exit state at (4,3).
- (b) The root-mean-square error in the estimate for  $U(1, 1)$ , averaged over 50 runs of 100 trials each.

# Learning from Rewards

## Temporal-difference learning

- use the observed transitions to adjust the utilities of the observed states so that they agree with the constraint equations.

### Temporal-difference (TD) equation

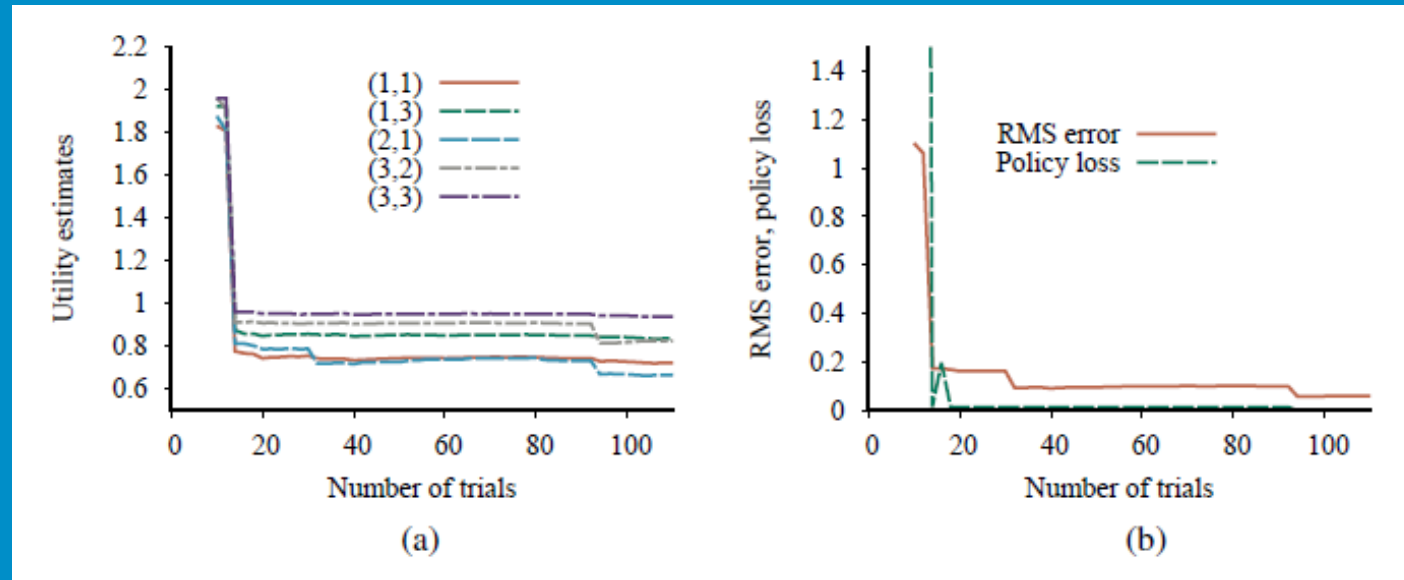
$$U^\pi(s) \leftarrow U^\pi(s) + \alpha[R(s, \pi(s), s') + \gamma U^\pi(s') - U^\pi(s)].$$

- $\alpha$  is the **learning rate** parameter.
- uses the difference in utilities between successive states (and thus successive times)
- adjusting the utility estimates toward the ideal equilibrium that holds locally when the utility estimates are correct
- does not need a transition model to perform its updates.

# Learning from Rewards

TD adjusts a state to agree with its <i>observed</i> successor	ADP adjusts the state to agree with all of the successors that might occur, weighted by their probabilities
TD makes a single adjustment per observed transition	ADP makes as many as it needs to restore consistency between the utility estimates $U$ and the transition model $P$ .
For each observed transition, the TD agent can generate a large number of imaginary transitions	can generate more efficient versions of ADP by directly approximating the algorithms for value iteration or policy iteration.

# Learning from Rewards



Performance of the exploratory ADP agent using  $R^+ = 2$  and  $N_e = 5$ .

(a) Utility estimates for selected states over time.

(b) The RMS error in utility values and the associated policy loss.

# Active Reinforcement Learning

- Passive agent has a fixed policy that determines its behavior
- Active learning agent gets to decide

## Exploration

- Refers to the agent's strategy of seeking information about the environment by trying different actions.
- A purely greedy agent sticks to its policy and does not explore alternative actions, potentially missing the optimal route.
  - Also termed as a greedy agent: greedily takes the action that it believes
  - Sometimes this approach pays off and sometimes it does not
  - Overlook that actions do more than provide rewards (Actions provide information in the form of percepts in the resulting states)
- Not greedy in immediate next move but **Greedy in Limit of infinite exploration (GLIE)**
  - GLIE scheme must try each action in each state an unbounded number of times
  - to avoid having a finite probability that an optimal action is missed.



# Active Reinforcement Learning

## Safe Exploration

- Many actions are **irreversible**
  - no subsequent sequence of actions can restore the state to what it was before the irreversible action was taken
  - Worst case: agent enters **absorbing state** (no actions have any effect/rewards)
- choose a policy that works reasonably well for the whole range of models that have a reasonable chance of being the true model
  - even if the policy happens to be suboptimal for the maximum-likelihood model.
- **Three mathematical approaches:**
  - i) Bayesian reinforcement learning
  - ii) Exploration POMDP
  - iii) Robust control theory

# Active Reinforcement Learning

## Temporal-difference Q-learning

- **Q-learning** method avoids the need for a model by learning an action-utility function  $Q(s, a)$  instead of a utility function  $U(s)$ .
- derive a model-free TD update for the Q-values

$$Q(s, a) \leftarrow Q(s, a) + \alpha [R(s, a, s') + \gamma \max_{a'} Q(s', a') - Q(s, a)] .$$

- No transition model
- **SARSA** (for state, action, reward, state, action).
  - updates with the Q-value of the action  $a'$  that is actually taken:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [R(s, a, s') + \gamma Q(s', a') - Q(s, a)] ,$$

# Active Reinforcement Learning

```
function Q-LEARNING-AGENT(percept) returns an action
  inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r$ 
  persistent:  $Q$ , a table of action values indexed by state and action, initially zero
                $N_{sa}$ , a table of frequencies for state–action pairs, initially zero
                $s, a$ , the previous state and action, initially null

  if  $s$  is not null then
    increment  $N_{sa}[s, a]$ 
     $Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
     $s, a \leftarrow s', \operatorname{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a'])$ 
  return  $a$ 
```

An exploratory Q-learning agent. It is an active learner that learns the value  $Q(s, a)$  of each action in each situation. It uses the same exploration function  $f$  as the exploratory ADP agent, but avoids having to learn the transition model.

# Applications of Reinforcement Learning

## Applications in game playing

- deep Q-network (DQN) system, the first modern deep RL system by DeepMind
- DQN was trained separately on each of 49 different Atari video games
- learned to drive simulated race cars, shoot alien spaceships, and bounce balls with paddles
- DeepMind's ALPHAGO beat the best human players

## Application to robot control

- cart–pole balancing problem, also known as the inverted pendulum
- radio-controlled helicopter flight
  - used policy search over large MDPs
  - often combined with imitation learning and inverse RL given observations of a human expert pilot

# Example1 - Policy Based agent

- Pong:
  - <https://www.youtube.com/watch?v=YOW8m2YGtRg>
  - Policy function:
    - Input: pixel intensities
    - Output: velocity of paddle

# Example2 - Q-learning

- Google DeepMind's Deep Q-learning playing Atari Breakout
  - <https://www.youtube.com/watch?v=V1eYniJ0Rnk>
  - Model Input:
    - Pixel intensities recorded over the last 4 time steps.
  - Model Output:
    - Anticipated future reward for each potential action.

# Applications of Reinforcement Learning



- (a) Setup for the problem of balancing a long pole on top of a moving cart. The cart can be jerked left or right by a controller that observes the cart's position  $x$  and velocity  $x'$ , as well as the pole's angle  $\vartheta$  and rate of change of angle  $\vartheta'$ .
- (b) Six superimposed time-lapse images of a single autonomous helicopter performing a very difficult "nose-in circle" maneuver. The helicopter is under the control of a policy developed by the PEGASUS policy-search algorithm (Ng *et al.*, 2003). A simulator model was developed by observing the effects of various control manipulations on the real helicopter; then the algorithm was run on the simulator model overnight. A variety of controllers were developed for different maneuvers. In all cases, performance far exceeded that of an expert human pilot using remote control. (Image courtesy of Andrew Ng.)

# Summary

- A **model-based reinforcement learning** has a transition model  $P(s^t / s, a)$  for the environment and learns a utility function  $U(s)$ .
- A **model-free reinforcement learning** agent may learn an action-utility function  $Q(s, a)$  or a policy  $\pi(s)$ .
- Utility learning approach:
  - Direct utility estimation
  - Adaptive dynamic programming (ADP)
  - Temporal-difference (TD)
- Reward shaping and hierarchical reinforcement learning are helpful for learning complex behaviors
- Policy-search methods operate directly on a representation of the policy, attempting to improve it based on observed performance
- Apprenticeship learning through observation of expert behavior can be an effective solution when a correct reward function is hard to specify.